RECIPE: Relaxed Sound Predictive Analysis

Abstract—The abstract goes here.

I. INTRODUCTION

In a multithreaded application, a data race occurs when two concurrent threads access the same memory location, such that (i) at least one of the accesses is write access, and (ii) no explicit mechanism is enforced to prevent the threads from simultaneous access to the location [8]. Race detection is important not only because races often reveal bugs in the implementation, but also as the basis of other techniques and analyses like atomicity checking [6], [7], [10], data-flow analysis [1] and record/replay [2], [5].

In this paper, we focus on *sound* race detection, whereby reported races are guaranteed to be real. This is the key requirement for adoption of the tool by developers [?]. Ensuring soundness is a difficult challenge, which mandates dynamic forms of analysis. Indeed, extensive research has been carried out on dynamic race detection. The general goal has been to derive constraints from a given execution trace on event reordering, and check for the remaining reorderings whether they disclose data races. The key question then becomes about constraint extraction.

Existing Approaches: Numerous techniques have been proposed to date for race detection, which either sacrifice soundness [4], [?], [?], [?] or have significant coverage limitations [?], [?], [?]. Recently *predictive analysis* has emerged as a promising alternative, whereby a single trace is considered, and races are discovered by permuting the execution schedule governing the trace [?], [?].

Predictive analysis guarantees soundness, and has the potential for high coverage, though current predictive analyses suffer from two major limitations:

1) All the permutations attempted by the analysis must preserve the flow dependencies exhibited by the original trace. As an example, if in the original trace thread t_1 reads a shared variable that was written by thread t_2, then the same must hold in the permuted trace.
2) The analysis cannot step outside the boundaries of the input trace, e.g. by exploring branches that were not followed in that trace.

Both of these constraints are a conservative means of ensuring soundness. In an empirical study we conducted, which we describe in Section VI, we found that these restrictions often lead to serious loss in coverage.

As an illustration, we refer to the trace in Figure 1, where the statements in red denote an unexplored branch outside the trace. Existing predictive analyses are unable to detect the race between lines 2 and 5 because of the dependence between lines 3 and 4, which is a barrier to the needed reorderings.

```
x = 0; y = 0; z = 3;
          z = 2;
─────────────────────────────
    T₁              T₂
1: y = 3;
2: x = 1;
3: y = 5;
              4: if (y > z)
              5:  print(1/x);
              6: else
              7:  print(2/x);
```

Fig. 1. Example illustrating ordering constraints beyond synchronization primitives

A second race that is missed, as it involves the statements in red, is between lines 2 and 7.

Our Approach: We describe a novel approach to predictive race detection, which we have implemented as the RECIPE analysis tool, that is able to relax both of these constraints. Compared to the state of the art [?], RECIPE is able to detect x1.5 more races with only the first restriction relaxed, and x2.5 more races with both restrictions relaxed.

The departure point of our approach is to consider the explicit values of shared memory locations instead of dependencies between trace events. The gain in coverage is twofold:

1) Reorderings that violate the original dependence structure but preserve the branching history become possible. The values read by a branching statement may change so long as the branch condition evaluates to its original truth value.
2) In addition, in certain cases we can guarantee soundness while directing execution toward branches that diverge from the original trace. This mandates that the effects of an unexplored branch can be modeled precisely, which holds frequently in practice because the original and new branchees often access the same set of shared memory locations.

Returning to the example in Figure 1, we demonstrate how value-based reasoning improves coverage. First, value-based reasoning reveals that trace [1, 4, 5, 2, 3] is feasible, since the first assignment, y = 3, satisfies the condition y > z. Hence the race is discovered. A second race, between lines 2 and 7, is detected by negating the condition. This is possible because the effects of the else branch can be modeled precisely. Negation leads to an execution starting at line 4 (where y ≡ 0). The race between lines 2 and 7 then becomes visible.

Concretely, we achieve value-level (rather than dependence-level) granularity computationally via a unique encoding of the input execution trace as a constraint system. The constraints

```
s ::=
                            y = x.f |    (heapr)
                            x.f = y |    (heapw)
                        z = x bop y |    (assign)
                 if (x bop y) goto ... |  (branch)
                    lock(l) | unlock(l) |    (sync)
      fork(t) | join(t)| begin(t) | end(t)|    (thread)
```

TABLE I
LANGUAGE SYNTAX

are then processed by a satisfiability checker. This gives us the flexibility to explore nontrivial constraints, such as forcing a context switch after the first assignment to y or negating the condition in Figure 1. Importantly, these manipulations are beyond what dependence-based reasoning can achieve, since this view of the trace is too conservative.

Contributions: This paper makes the following principal contributions:

II. TECHNICAL OVERVIEW

In this section, we walk the reader through a detailed technical description of our approach based on the example in Figure 1. As input, we assume (i) a program P as well as (ii) a trace of P recorded during a dynamic execution.

A. Preliminaries

To facilitate our presentation, we first introduce some terminologies used throughout this paper.

Trace A trace τ is a sequence of events recorded during the observation run.

Event An event, $e = < t, id, inst, map >$, is a concrete representation that captures the details about the runtime execution of a static instruction $inst$, where t denotes the thread that issues the event, id is the id associated with each event, $inst$ is the static instruction and map maps the variable in the instruction to the runtime values. The key property of id is *uniqueness*, i.e., any two events in the trace own different ids. A natural id for each event is its index in the trace. We denote the thread and id as t^e and i^e respectively.

The instructions are three-address instructions involving at most three operands, which modern compilers commonly support. Specifically, we are interested in the following types of instructions.

When the variable does not appear on the left hand of an equation, such as y in $x.f = y$, it may refer to a variable, a constant or event object creation expression $new(...)$. The bop stands for the binary operator, which may refer to $+, -, *, /, \%, \wedge, \vee$ in the assignment, or refer to $<, >, =$ in the branch. The listed instructions are so basic that it can capture most events in a common imperative programming language. For example, the method call can be modeled by the above list, where the argument passing and value return are modeled as assignment events, and the virtual function resolution is modeled as the branch event conditioned on the type of the caller object.

The map stores the value for each variable observed in the dynamic run. For simplicity, we use the notion $[\![x]\!]^e$ to retrieve the value associated with the variable x in e.

In addition, we propose the projection operation $|$ to help the reasoning of the trace: $\tau|t$ contains only the (ordered) events from the thread t; $\tau|i$ contains only the event of which the id equals i; $\tau| \leq e$ denotes the prefix of e, i.e., the events preceding the event e; $\tau| \geq e$ denotes the events after e; $\tau| >= e_1 \wedge <= e_2$ denotes the events between e_1 and e_2.

B. Constraint System

We begin with an explanation of the encoding process. The resulting formula is provided in Figure 2. We describe the conjuncts comprising the formula one by one.

Program Order: The first set of constraints reflects ordering constraints between program statements. We use the symbol O_i to denote the ith program statement, which yields the following formula for Figure 1:

$$O_1 < O_2 < O_3 \wedge O_4 < O_5 \wedge O_4 < O_7$$

That is, the first 3 statements are totally ordered, and the if statement executes before the body (either O_5 or O_7).

Variable Definitions: The next set of constraints, denoted $z^i = k$, capture variable definitions: Variable z is assigned value k at statement i. For our running example, we obtain:

$$x^0 = 0 \wedge y^0 = 0 \wedge y^1 = 3 \wedge x^2 = 1 \wedge y^3 = 5$$

As an example, $y^3 = 5$ denotes that the value assigned to variable y at line 3 is 5.

Thread Interference: To express inter-thread flow constraints, we utilize expressions of the form $R_z^i = z^j$, which denotes that line i reads variable z, and the definition it reads comes from line j. The resulting formula for our example is

$$
\begin{aligned}
& (R_y^4 = y^0 \wedge O_4 < O_1) \\
\vee \; & (R_y^4 = y^1 \wedge O_1 < O_4 < O_3) \\
\vee \; & (R_y^4 = y^3 \wedge O_3 < O_4)
\end{aligned}
$$

Notice, importantly, that the formula combines flow constraints with order constraints, which are essential to determinize the value read at a given statement. As an example, $(R_y^4 = y^1 \wedge O_1 < O_4 < O_3)$ means that the value of y read at line 4 is that set at line 1 assuming an execution order whereby the first statement executes followed by the fourth then third statements.

Path Conditions: To preserve path conditions while potentially permitting dependence-violating reordering (e.g., a context switch after line 1 in Figure 1), we model explicitly the condition. For the running example (ignoring the statements in red), this yields:

$$R_y^4 > 2$$

That is, the value of variable y read at line 4 is greater than 2. Indeed, this constraint is satisfied by the assignments to y both at line 1 and at line 3.

$$
\begin{array}{lll}
& (O_1 < O_2 < O_3 \wedge O_4 < O_5 \wedge O_4 < O_7) & \textbf{(program order)} \\
\wedge & (\mathbf{x}^0 = 0 \wedge \mathbf{y}^0 = 0 \wedge \mathbf{y}^1 = 3 \wedge \mathbf{x}^2 = 1 \wedge \mathbf{y}^3 = 5) & \textbf{(variable definitions)} \\
\wedge & ((R_\mathbf{y}^4 = \mathbf{y}^0 \wedge O_4 < O_1) \bigvee (R_\mathbf{y}^4 = \mathbf{y}^1 \wedge O_1 < O_4 < O_3) \bigvee (R_\mathbf{y}^4 = \mathbf{y}^3 \wedge O_3 < O_4)) & \textbf{(thread interference)} \\
\wedge & R_\mathbf{y}^4 > 2 & \textbf{(path conditions)} \\
\wedge & O_2 = O_5 & \textbf{(race condition)}
\end{array}
$$

Fig. 2. RECIPE encoding of the trace in Figure 1 as a constraint system

Race Condition: The final constraint, forcing the check whether a particular race is feasible, is to demand that two conflicting statements occur at the same time. For the potential race between lines 2 and 5, we obtain:

$$O_2 = O_5$$

This asserts that both statements occur simultaneously, which — together with the other constraints — guarantees the feasibility of the predicted race if a solution is found for the overall constraint system.

Unexplored Branches: Beyond the encoding steps so far, which focus on the given trace, we can often encode constraints along unexplored branches. In our running example, this is essential to discover the race between lines 2 and 7.

The conflicting accesses in this case, determined based on static analysis of P, are expressed as $O_2 = O_7$. In addition, we negate the path condition, thereby obtaining $R_\mathbf{y}^4 \leq 2$ in place of $R_\mathbf{y}^4 \leq 2$.

Constraint Solving: Having conjoined the formulas from the different encoding steps into (i) a global representation of all feasibility constraints (path, ordering, assignment and other constraints) and (ii) the requirement for a given race to occur (expressed as simultaneous execution of the conflicting accesses), we discharge the resulting formula to an off-the-shelf constraint solver, such as Z3 or Yices. If successful, the solver returns a solution for the specified constraints.

In particular, the solution discloses a feasible trace that gives rise to the race at hand. The trace is identified uniquely via the order enforced in the solution over the variables O_i, which represent scheduling order.

III. RELAXATION OF FLOW DEPENDENCIES

We begin by describing the first relaxation, whereby the dependence structure defined by the original trace is potentially violated while retaining soundness.

A. Trace Encoding

The first step of the process is to rewrite the concrete execution trace into a symbolic form that enables RECIPE to derive reordering constraints. We describe how RECIPE encodes a trace into symbolic form in two steps. We first start with "local" intraprocedural events — i.e., transitions that manipulate memory accessed by at most a single thread during an invocation-free run — and then extend our encoding scheme to the entire set of possible events.

Basic Encoding: Local Accesses: The fundamental encoding transformation is to induce Static Single Assignment (SSA) form on the raw trace, such that a variable is defined exactly once. In this way, def/use chains become explicit, and encoding of trace events as constraints is simplified. As an illustration, trace

```
1: x=1; 2: x<3; 3: x=3;
4: y=1;
5: z=x+y
```

becomes

```
1: x¹=1; 2: x¹<3; 3: x³=3;
4: y⁴=1;
5: z⁵=x³+y⁴
```

(For readability, we version variables according to the line number of their definition.)

Beyond the default SSA rewriting algorithm, we apply a specialized transformation to handle local heap accesses. While in general determining whether a given heap access is local is an undecidable problem, in the context of a concrete execution trace this determination is straightforward. We can thus improve upon the baseline SSA form, where def/use chains over heap accesses are ignored. Indeed, to ensure feasibility, we must account for such accesses. Hence, we replace accesses to field $\mathbf{f}$ of local object o with a fresh local variable $\mathbf{l}_{o.\mathbf{f}}$, where $o.\mathbf{f}$ denotes the memory location itself and not its value. This is done prior to the standard SSA transformation, and under the assumption that in the predicted run, the base object o remains the same.

Full Encoding: Method Calls and Shared Accesses: To complete the encoding algorithm, we next explain how shared heap accesses and method calls are dealt with. Similarly to the case of local heap accesses, we take advantage of the fact that the full execution trace is available for analysis to detect shared accesses (the dual of local accesses) as well as call-site resolutions (which is more relevant in practice, where we support virtual method calls).

To account for method invocations, we induce additional context on variable accesses. Instead of recording only the version of a variable, we also represent as part of the variable's identifier its enclosing method and its invocation counter. Thus, symbol $\mathbf{m}^2 \colon \mathbf{v}^7$ is interpreted as the definition of local variable $\mathbf{v}$ within the second invocation of method $\mathbf{m}$ that occurs at trace index 7.

In this way, the trace is flattened. Method invocations are substituted with qualified variable names, and the rest of the encoding steps, described above, remain unchanged. Because

the trace is finite and fully resolved, challenges such as looping, recursion and mutual recursion are all obviated.

The final aspect of our encoding process is representation of shared heap accesses. As illustrated in Section II, this is done via designated symbols:

- $W_{\mathtt{v}}^{k}$ denotes write access to variable $\mathtt{v}$ at trace index k (where $\mathtt{v}$ is assumed to be a qualified identifier to account for the interprocedural setting).
- Analogously, $R_{\mathtt{v}}^{k}$ denotes that $\mathtt{v}$ is read at index k.

The main motivation to encode shared heap accesses using special symbols is to simplify downstream processing, and in particular, the definition of candidate races. We could encode the question of shared versus local accesses as additional constraints (requiring that a variable be accessed by more than one thread), but that would be more complicated and less efficient.

B. Constraint System

Having explained how the trace is encoded, we now describe in detail how constraints are derived from the trace, such that any permutation considered by the analysis is guaranteed to represent a feasible execution schedule.

Intra-thread Order: The first set of constraints reflects control flow within the individual threads, which must remain unchanged under the reordering transformation. Given input trace t, this is expressed as the following formula:

$$\forall \tau, \tau' \in t. \quad \mathsf{thread}\ \tau \equiv \mathsf{thread}\ \tau'. \\ \mathsf{index}\ \tau < \mathsf{index}\ \tau' \Rightarrow O_\tau < O_{\tau'}$$

The logical variables O_x express ordering constraints. The requirement, as stated above, is that these variables reflect the same order as the projection of index onto individual threads.

Race Condition: Given pair τ and τ' of events that both access a common memory location ℓ, we demand that

$$\mathsf{thread}\ \tau \neq \mathsf{thread}\ \tau' \\ \wedge \quad (\ell \in \mathsf{writeset}\ \tau \vee \ell \in \mathsf{writeset}\ \tau') \\ \wedge \quad O_\tau = O_{\tau'}$$

That is, (i) events τ and τ' are executed by different threads, (ii) at least one of the events performs write access to ℓ (as judged by the writeset membership check), and (iii) the events occur simultanesouly. This is a direct logical encoding of the definition of a race condition.

Path Constraints: The requirement with respect to branching is that the prefix of the original trace t up to the pair τ and τ' of candidate racing events remains identical in the predicted trace t'. More accurately, under the assumption that $\mathsf{index}\ \tau < \mathsf{index}\ \tau'$, we require that

$$\bigwedge_{\tau'' \in t \cap \mathbf{bexp}.\ \tau'' \in \mathsf{pre}\ \tau'} [\![\mathsf{stmt}\ \tau'']\!]\ t \equiv [\![\mathsf{stmt}\ \tau'']\!]\ t'$$

where stmt is a helper function that obtains the code statement incident in a given transition. That is, all branching transitions up to τ' (which occurs after τ) preserve their boolean interpretation under t'. This ensures that there are no divergences from the path containing the racing events, though beyond

that path any feasible continuation is permitted. Importantly, contrary to [**?**], we do not pose the requirement that the values flowing into branching statements remain the same, but suffice with the relaxed requirement that the evaluation of branching expressions is invariant under the input and predicted traces.

Variable Definitions: A final requirement for the basic setting is that left-hand variables are defined according to the same right-hand variables as before. That is, version i of variable $\mathtt{u}$ is defined as version j of variable $\mathtt{v}$ in input trace t, then the same remains true in predicted trace t'. This is enforced as the formula

$$\bigwedge_{\tau'' \in t \cap \mathbf{asgn}.\ \tau'' \in \mathsf{pre}\ \tau'} \mathsf{stmt}\ \tau'' \in t'$$

where we again assume that $\mathsf{index}\ \tau < \mathsf{index}\ \tau'$. That is, the same statement occurring in t is also present in t' (though the transitions may differ). Since the statements of t' are a permutation of the statements of t, we are assured that use/def flow is constrained appropriately.

Relaxation of Flow Dependencies: We now move to the novel feature of RECIPE, which is its ability to explore execution schedules that depart from the data flow exhibited in the original trace. More precisely, RECIPE is able to relax flow dependencies in the original trace, whereby a thread reads a shared memory location written by another thread, while enforcing feasibility. This is strictly beyond the coverage potential of existing predictive analyses, which restrict trace transformations to ones where any read access to a shared memory location must correspond to the same write access as in the original trace.

To ensure feasibility under relaxation of flow dependencies, we need to secure the link between the execution schedule and the write/read flow. As an illustration from the example in Figure 1, $R_{\mathtt{y}}^{4} = \mathtt{y}^1 \wedge O_1 < O_4 < O_3$ specifies that in a schedule where thread T_1 executes line 1, then T_2 executes line 4, and then the schedule switches again to T_1 to execute line 3, the read access to $\mathtt{y}$ at line 4 obtains the value assigned to $\mathtt{y}$ at line 1: $R_{\mathtt{y}}^{4} = \mathtt{y}^1$.

The full and general constraint formula, given read R_ℓ of location ℓ as part of event e with set $\mathcal{W}$ of matching write events (i.e., events including write access to ℓ), takes the following form:

$$\bigvee_{e_w \in \mathcal{W}} \begin{array}{l} (R_\ell = \ell^{\mathsf{index}\ e_w}) \\ \wedge \quad O(e_w) < O(e) \\ \bigwedge_{e' \in \mathcal{W} \setminus \{e_w\}} (O(e') < O(e_w) \vee O(e) < O(e')) \end{array}$$

This disjunctive formula iterates over all matching write events, and demands for each that (i) it occurs prior to the read event ($O(e_w) < O(e)$) and (ii) all other write events either occur before ($O(e') < O(e_w)$) it or after the read event ($O(e) < O(e')$).

An important concern that arises due to relaxation of flow dependencies is that heap accesses may change their meaning. As an illustration, we refer to Figure 3. While the read at line 7 appears to match the write at line 5, this is conditioned on the read at line 6 being linked to the assignment at line

	T_1	T_2
1:	x1 = new();	
2:	x2 = new();	
3:	y = x1;	
4:	y = x2;	
5:	x2.f = 5;	
6:		z = y;
7:		w = z.f;

Fig. 3. Example illustrating the need to account for heap accesses during trace transformation

4. If the predicted run violates this link, then feasibility is no longer guaranteed. In particular, if reordering results in z being assigned the first rather than second allocated object, then the write at line 5 no longer matches the read at line 7.

To address this challenge, we enhance the constraint system with the requirement that heap objects that are dereferenced before the candidate racing events retain their original address in the predicted trace. This achieves three guarantees: First, matching events in the original trace are guaranteed to also match in the predicted trace. Second, candidate races in the original trace remain viable in the predicted trace. Finally, in a practical setting involving virtual method calls (which goes beyond our core grammar in Table I), call-site resolutions are the same across the original and predicted traces. Notice that in this setting, we consider as relevant not only the targets of field dereferences but also the targets of virtual method invocations.

Formally, given pair τ and τ' of candidate racing events such that index $\tau <$ index τ', we require that

$$\bigwedge_{\tau'' \equiv \text{y=x.f} \in t \cap \textbf{heapr}. \ \tau'' \in \text{pre } \tau'} \text{env } \sigma(t, \tau'') \text{ x} = \text{env } \sigma(t', \tau'') \text{ x}$$

where env is the state mapping from local variables to their value and $\sigma(t, \tau'')$ (resp. $\sigma(t', \tau'')$) is the state arising at trace t (resp. t') immediately before event τ''. This constraint fixes that all heap dereferences up to the later of the candidate racing events retain their original base object as in the predicted trace t'.

IV. EXPLORATION OF UNEXECUTED BRANCHES

We now switch to the second feature of RECIPE, which is its ability to reason about unexplored branches.

Relaxation of Branching Decisions: Beyond relaxing flow dependencies, RECIPE is also capable of exploring code branches that were not executed in the original trace. This is achieved via a symbolic representation of the input trace. Given branching event e_b,

1) model the branching condition symbolically to execute along the negation of the original branch;
2) apply depth-first search (DFS) to uncover all execution suffixes under the unexplored branch (which may itself contain branching statements); and
3) for each suffix t^s,

a) truncate the original trace at e_b yielding prefix t_s; and
b) concatenate t_s with the negation of e_b followed by t^s.

Constraint solving is applied to each of the resulting traces analogously to the original trace.

We emphasize that exploration of new execution paths is subject to all the known limitations of symbolic execution, including in particular loop structures and object allocation. RECIPE currently fails if (i) the unexplored branch containts loops or (ii) there are object references that cannot be fully resolved at the branching point. In case of failure, RECIPE moves on to other branches. We demonstrate in Section VI that despite these limitations, the increase in coverage thanks to exploration of new paths is significant.

V. SOUNDNESS PROOF

In this section, we prove formally the soundness of RECIPE.

Researchers [3], [9] propose two axioms, *prefix closedness* and *local determinism*, which we adapt to establish our feasibility guarantee. Suppose $\mathcal{F}$ denotes the domain of feasible traces.

Intuitively, Prefix closedness means that if a trace is feasible, then the trace is still feasible after we discard the events after a point. The underlying reason is that the events after a point cannot affect the events before it. More formally,

Axiom 1 (Prefix closedness). *If $t_1 \ t_2 \in \mathcal{F}$, then $t_1 \in \mathcal{F}$.*

Intuitively, local determinism says that only the previous events of the same thread determine the existence of an event. For example, the existence of an event is only determined by the branch thread-locally. However, the value read by an event may be affected by another thread because a shared read reads from the most recent write, which may come from a different thread. The value in other events, such as branch or write, are locally determined too. More formally,

Axiom 2 (Local determinism). *Assume $t_1 e_1, t_2 \in \mathcal{F}$, and $t_1|thread(e_1) \approx t_2|thread(e_1)$, where $\approx$ denotes the two traces are equal if the data values in the read and write events are ignored, Each event is determined by the previous events in the same thread. There are four cases:*

- **Branch** . *If $op(e_1)$=branch, and e_1 is in the form of $x < y$ (without loss of generality), then there exists values v_x, v_y such that $t_2 e_1[v_x/data_x, v_y/data_y] \in \mathcal{F}$. Here $e_1[v_x/data_x]$ represents that we replace the value of x from $data_x$ to v_x in e_1.*
- **Read** . *If $op(e_1)$=read, and e_2 is a read event that is identical to e_1 except that it may read a different value, and $t_2 e_2$ is consistent, then $t_2 e_2 \in \mathcal{F}$.*
- **Write** . *If $op(e_1)$=write, and there is a value v such that $t_2 e_1[v/data] \in \mathcal{F}$.*
- **Others** . *If $op(e_1)$ is of type different from above and $t_2 e_1$ is consistent, then $t_2 e_1 \in \mathcal{F}$.*

Here, the consistency is defined as follows. A trace is consistency iff it satisfies (1) read-write consistency, i.e., read

event should contain the value written by the most recent write event, and (2) the synchronization consistency, e.g., the lock acquire and release of a lock should not be interleaved by the lock operations of the same lock, the begin event of a thread should follow the fork event of the parent thread.

The local determinism resembles the version in [3], [9]. A crucial difference is that, for the branch event, we do not require all reads thread locally before it to read the identical values as in the original trace $t_1 e_1$. Instead, we allow the previous reads to read different values as long as the branch event is evaluated to the same boolean value, which guarantees the existence of the following events. Such relaxation allows us to find more races, while still preserving the feasibility.

An important assumption of ours is, we assume all local accesses are recorded in the trace so that we can determine the boolean status of a branch event computationally. [3], which assume the local accesses are not recorded and cannot determine the status computationally, have to enforce the reads before a branch to read the same values, leading to a weaker feasibility criterion. We believe our axiom captures the real-world scenario more faithfully. Another important assumption is for the read-write consistency, i.e., we assume the heap invariant, so that we know exactly which reads correspond to which writes in the predicted run.

They are called axioms because they are the rules that the sequentially consistent system [9] should follow. We refer the readers to the detailed discussion [9], [3].

Theorem 1 (Soundness). *All traces in the closure $closure(t, \Phi)$, which includes t and is closed under the derivations in Axioms 1 2, are feasible.*

Proof. We order the traces in $closure(t)$ as, $t_0 = t, t_1, t_2, t_3, \ldots t_n$. Clearly, t_0 is feasible. In the following, we prove by induction, i.e., assume t_n is feasible, then t_{n+1} which is derived from $S = \{t_0, \ldots t_n\}$ should be feasible. There are several possible derivations.

- if t_{n+1} is a prefix of $t' \in S$, then $t_{n+1} \in \mathcal{F}$ because of the prefix closedness. t_{n+1} shares the same mappings as t'.
- if t_{n+1} is derived from $t^1 e^1, t^2 \in \mathcal{F}$, we have $t^1|thread(e^1) \approx t^2|thread(e^1)$.
  - $op(e^1) = branch$ and $e^1 = x < y$. Our solver computes a mapping θ^d of all relevant variables such that $\theta(x < y) = \theta^d(x < y)$, then $t^2 e_1[\theta^d(x)/\theta(x), \theta^d(y)/\theta(y)] \in \mathcal{F}$. Note that the mapping computed by solver may be different from the mapping in the trace t^1.
  - $op(e^1) = read$. Our solver ensures the read-write consistency, i.e., $t^2 e^2$ is consistent, then $t^2 e^2 \in \mathcal{F}$. Note that e^2 is the same as e^1 except that it may read a different value.
  - $op(e^1) = write$. There exists a value v such that $t^2 e^2[v/data] \in \mathcal{F}$, where v is included in the mapping computed by the solver.
  - $op(e^1)$ is of other types. Our solver ensures the

consistency, i.e., $t^2 e^1$ is consistent, then $t^2 e^2 \in \mathcal{F}$. □

Theorem 2 (Maximality). $\forall t', s.t., t'|th \approx t|th\ for\ any\ thread\ th, if\ t' \notin closure(t, \Phi), t'\ is\ infeasible.$

Proof. We sketch the proof. All traces that satisfy $t'|th \approx t|th\ for\ any\ thread\ th$ □

VI. EVALUATION

Our evaluation focuses on the effectiveness and scalability of our approach. To measure the effectiveness, we compare with the predictive analysis, RV []. We choose RV for two main reasons: (1) RV is the only open source predictive analysis, (2) RV represents the state-of-the-art approach, which is theoretically proven to have higher detection capability than other approaches.

Evaluation Method We conduct our experiment on a large set of applications, which are also used to evaluate RV. Specifically, the set includes large applications such as Jigsaw, Xalan, Lusearch. For the large applications that use the reflection, we adopt Tamiflex [] to to support the static analysis, which replaces the reflection calls with the concrete method calls recorded in the observation run. We omit the benchmark eclipse because the current version of Tamiflex leads to abnormal execution after the instrumentation, which throws exception during starting the main service. By applying our tool to such abnormal execution, we identify only 3 races, similar to the report of RV []. Besides, the benchmark montecarlo requires the input, which we downloaded from internet and simplified. For the large applications, we use the most lightweight configuration if possible.

As the detection capability depends on the observed run, for fair comparison, we monitor the execution once by recording all necessary information required by both approaches, and then apply both techniques to the monitored run. Besides, our reported data for RV may be different from the original report for two reasons: (1) the different observed runs lead to different set of races, (2) the original implementation of RV contains a bug in identification of the branches, which leads to the misses of many branches and the incorrect reduction of dependences during the analysis. We confirmed the bug with the author and fixed the bug in our experiments [1].

Our experiments and measurements were all conducted on an x86 64 Thinkpad W530 worksta- tion with eight 2.30GHz Intel Core i7-3610QM processors, 16GB of RAM and 6M caches. The workstation runs version 12.04 of the Ubuntu Linux distribution, and has the Sun 64-Bit 1.6.0_26 Java virtual machine (JVM) installed [2].

A. Effectiveness

Table II shows the main results of our analysis, which includes four sections: Trace (details about the trace), Races

---

[1] We report the bug in details with the test case https://sites.google.com/site/recipe3141/

[2] Note that Sun JDK 1.7 does not support the transformation of dacapo applications with tamiflex.

(detected races), Difference (comparison with RV) and Running time (the time taken by the analysis).

The Trace section includes the number of threads (Th), the number of shared reads ($Reads$) and shared writes ($Writes$), the subset of shared reads that read the base object references ($Base$), where the base object references are references used as the base/target in the following field reference or the method invocation, the number of branches (Br), the synchronization events ($Sync$) and the total number of events ($Total$), which includes local accesses in addition to the aforementioned events. Specifically, for the branch, we report it in the form of A/B, where A refers to the number of branches used by our analysis and B refers to the number of branches used by RV. To ensure that the predicted run sees the same base object at each shared read/write, RV inserts the artificial branch immediately in front of each shared access (and array accesses). We do not use such artificial branches.

We make interesting observations about the Trace section. The non-local events, i.e., all the events listed in Table II, occupy around 30% of the total trace in the first 7 benchmarks, but occupy less than 1% of the trace in the rest benchmarks, which have relatively more complex logic. The reads of the shared base objects occupy a small portion (1/3-1/10) of the total shared reads. The rest shared reads read only the primitive values or the references that are not the base references, e.g., the references involved only in the nullness check. Besides, our analysis involves significantly less branch events as compared to the RV approach. The difference plays an important role in the detection, which we will explain shortly.

The $Races$ section shows the number of races detected by Recipe-s, i.e., Recipe without exploring un-executed paths, Recipe, i.e., the fully-fledged version, and RV. By comparing the Recipe-s and Recipe, we find Recipe finds 100+ more races, which demonstrates the strength of exploring un-executed paths. Intuitively, Recipe-s predicts based on a single trace, while Recipe predicts based on multiple traces containing different execution paths. We also compare the Recipe version with RV version, as illustrated in the $Difference$ section, where $Diff$ shows the races found by Recipe but missed by RV, $Diff'$ shows the races found by RV but missed by Recipe.

First of all, Recipe finds many 150+ more races than RV. The reasons are multifold: (1) Recipe can reason about the accesses in the un-executed paths, while RV and Recipe-s can only reason about the accesses in the executed paths. (2) Recipe or Recipe-s allows the relaxation of the scheduling even if it breaks the inter-thread read/write dependence in the observed run. Recipe allows the majority of the shared reads, i.e., the reads of primitive values, to freely read from a different value from a different write as long as the value leads to the same branch decisions. RV, however, requires them to read the same values as in the observed run. (3) An critical optimization proposed by RV is to preserve dependences only for the reads before the preceding branches of the racy events, rather than all the reads. However, this optimization is underplayed by the fact that RV introduces huge amount of

artificial branches, i.e., one before each shared field access, which ensures the use of the same base objects. We get rid of such artificial branches and instead, rely on the small amount of base read events to ensure the use of the same base objects (Section **??**). Our strategy reduces the number of branches greatly and amplifies the effectiveness of the optimization. We also conduct case studies (Section **??**) to better illustrate the scenarios.

Another interesting observation is, although Recipe should produce all races found by RV in theory, Recipe may miss races found by RV in practice (Column $Diff'$), i.e., Recipe is not strictly more effective than RV in practice. The underlying reason is due to the limits of the constraint solver: (1) the solver cannot compute constraints with very complex arithmetic operations (2) the solver does not support some program constants such as the scientific notation, $3E - 10$.

The last section, Running time, compares the analysis time for both approaches. We find our approach is significantly slower than RV, e.g., RV often finishes within 200 seconds, while our approach may take more than 1 hour. This is because our approach needs to reason about the computation among variables inside the local access events, while RV needs to only reason about the order relations among the events.

Summary of advantage and weakness In general, our approach is flexible, which allows the relaxation of schedules and paths based on the fine-grained reasoning of value flows. As a result, it detects many more races compared to existing approaches. On the other hand, it is heavy weighted. The suggestion is to combine it with the lightweight approach with soundness guarantee such as RV, by treating RV as the preprocessing and instructing Recipe to skip those confirmed by RV. In this way, we can complement Recipe by finding those missed by it due to the limit of the constraint solver.

B. Case studies

We manually inspect the reported races in small applications to gain better understanding about our approach.

Relaxing the Inter-thread dependence Figure 4 demonstrates the relaxation of inter-thread dependence enabled by our approach. The code is from the benchmark bbuffer, where the line number is marked. Huang et al. [] detects the race between line 291 and line 400, but fails to detect the race between line 294 and line 400. The reason is as follows. In the observation run, the execution follows the order, lines 400, 291 and 294. For the event at line 294, its preceding branch at line 291 reads from line 400. Therefore, Huang et al. [] requires the predicted run to preserve the dependence between line 291 and line 400 so that line 291 reads exactly the same value and the branch takes the same branch decision. The dependence enforces the order constraint, $400 \rightarrow 291$, which further enforces the order $400 \rightarrow 291 \rightarrow 294$. Our approach does not require the existence of such dependence. Specifically, we allow line 291 to happen before line 400 in the predicted run as long as the value read by it leads to the same branch decision, which is true in this case. As a result, there is no order constraint between line 294 and line 400, and

TABLE II
RELAXED ANALYSIS

Benchmarks	Trace							Races			Difference		Running time (sec)	
	Th	Reads	Writes	Base	Br	sync	Total	Recipe-s	Recipe	RV	Diff	Diff'	Recipe	RV
critical	3	13	7	5	2/13	6	78	8	8	8	0	0	8	2
airline	11	45	15	4	32/82	30	317	9	9	9	0	0	490	4
account	3	46	21	10	3/47	6	227	2	5	5	3	3	41	4
pingpong	4	7	7	3	0/15	6	111	1	1	1	0	0	19	1
bbuffer	4	640	118	10	634/1.1K	217	3.3K	13	25	9	21	5	62	5
bubblesort	26	1.3k	966	121	155/2.8K	322	8.4K	7	7	7	0	0	3295	3
bufwriter	5	165	52	75	16/130	44	525	4	10	2	8	0	63	9
mergesort	5	38	33	5	15/472	28	1.7K	3	10	3	10	0	37	5
raytracer	2	31	5	12	314/8,2K	676	94.5K	4	6	4	2	0	47	2
montecarlo	2	5	86	2	1.9K/38.2K	21.1K	1.9M	1	4	1	3	0	1	17
moldyn	2	605	61	104	19.6K/52.6K	62	203.4K	6	14	2	12	0	2842	1
ftpserver	28	684	299	71	4.4K/233.3K	78.2K	3.9M	99	152	57	108	13	811	153
jigsaw	12	525	702	211	63.2K/467.9K	86.7K	5.5M	17	23	8	15	0	33	7
sunflow	9	2.1K	1.3K	473	201.3K/827.0K	50K	7.1M	38	78	20	69	11	4520	22
xalan	9	1.4K	0.9K	209	15.7K/103.2K	190.1K	6.6M	2	6	2	4	0	5317	10
lusearch	10	2.3K	0.5K	715	22.2K/164K	93.2K	9.1M	27	49	14	38	5	5430	8

the two forms a race. We re-replay such race easily using the eclipse IDE breakpoints.


Fig. 4. Relaxation of Inter-thread dependence

Relaxing the Paths Figure 5 demonstrates how our approach relaxes the schedulings to account for the unexecuted paths. Each thread invokes the method Sorting, which recursively starts two children threads if there are two more available entries in the thread pool (lines 8-11), or starts one child thread if there is only one available entry (lines 4-5). The availability is computed through the static method available with the constant `total`, which is equal to 5. The shared variable `alive` denotes the used entries.

Initially there is one sorting thread. After it starts Thread 1 and Thread 2, there are three threads alive and only two more entries are available. Suppose in the observation run, Thread 2 consumes both thread entries and starts the children Thread 3 and Thread 4 (not shown), updating `alive` to 5, then Thread 1 cannot execute the branch at lines 4-5. Huang et al. [] require the predicted run to preserve the dependence denoted by the dotted line since the dependence affects the branch condition at line 2. As a result, the predicted run follows the same branch decision and cannot reason about the code at lines 4-5. Our analysis does not have such limitation.

Instead, it allows the predicted run to reason about the unexplored code. Specifically, it does not need to preserve the dependence from line 11 to line 1 and it allows line 1 (Thread 1) to read from line 9 (Thread 2). As a result, the branch condition guarding the unexplored branch is evaluated to be true, enabling the unexplored path in the constraint solver. Finally, the solver identifies the race between line 5 (Thread 1) and line 1 (Thread 3). Note that the two lines are synchronized on different locks [3].

VII. VALIDITY THREAT

Java method can have 65535 bytecode instructions maximally. Therefore, we count the number of bytecode instructions inside each method, if the number exceeds 65525, we avoid instrumenting the method. The consequence is that, we will miss the races inside the method. In this case, we specify the variables read from the method to be equal to their concrete values in the constraints. THe constraint solver can proceed safely without being affected by such methods.

We do not support the boolean operations such as &, bit operators <<, which contributes to most of our misses.

VIII. CONCLUSION

The conclusion goes here.

ACKNOWLEDGMENT

The authors would like to thank...

REFERENCES

[1] R. Chugh, J. W. Voung, R. Jhala, and S. Lerner. Dataflow analysis for concurrent programs using datarace detection. In *Proceedings of the 2008 ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '08, pages 316–326, New York, NY, USA, 2008. ACM.

[3]We abbreviate the synchronized keyword as sync.

```
static sync available(){   return total−alive; }
```

Thread 1

Thread 2

```
Sorting(...)
{

    // child1=...
1   y= available();
2   if(y==0){···}
3   else if(y==1){
4       child1.start();
5       sync(this){alive++;}
6   }
7   else{
8       child1.start();
9       sync(this){alive++;}
10      child2.start();
11      sync(this){alive++;} ...
    }


}
```

```
...
9   sync(this){alive++;}
...
11  sync(this){alive++;}
```

Thread 3

```
1   available();
```

Fig. 5. Relaxation of Paths

[2] J. Huang, P. Liu, and C. Zhang. Leap: Lightweight deterministic multi-processor replay of concurrent java programs. In *Proceedings of the Eighteenth ACM SIGSOFT International Symposium on Foundations of Software Engineering*, FSE '10, pages 207–216, New York, NY, USA, 2010. ACM.

[3] J. Huang, P. O. Meredith, and G. Rosu. Maximal sound predictive race detection with control flow abstraction. In *PLDI*, 2014.

[4] V. Kahlon and C. Wang. Universal causality graphs: A precise happens-before model for detecting bugs in concurrent programs. In *Proceedings of the 22Nd International Conference on Computer Aided Verification*, CAV'10, pages 434–449, Berlin, Heidelberg, 2010. Springer-Verlag.

[5] D. Lee, P. M. Chen, J. Flinn, and S. Narayanasamy. Chimera: Hybrid program analysis for determinism. In *Proceedings of the 33rd ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '12, pages 463–474, New York, NY, USA, 2012. ACM.

[6] S. Lu, S. Park, E. Seo, and Y. Zhou. Learning from mistakes: A comprehensive study on real world concurrency bug characteristics. In *Proceedings of the 13th International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS XIII, pages 329–339, New York, NY, USA, 2008. ACM.

[7] S. Lu, J. Tucek, F. Qin, and Y. Zhou. Avio: Detecting atomicity violations via access interleaving invariants. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS XII, pages 37–48, New York, NY, USA, 2006. ACM.

[8] S. Savage, M. Burrows, G. Nelson, P. Sobalvarro, and T. Anderson. Eraser: A dynamic data race detector for multi-threaded programs. In *Proceedings of the Sixteenth ACM Symposium on Operating Systems Principles*, SOSP '97, pages 27–37, New York, NY, USA, 1997. ACM.

[9] T. F. Serbanuta, F. Chen, and G. Rosu. Maximal causal models for sequentially consistent systems. In *RV*, 2012.

[10] Y. Shi, S. Park, Z. Yin, S. Lu, Y. Zhou, W. Chen, and W. Zheng. Do i use the wrong definition?: Defuse: Definition-use invariants for detecting concurrency and sequential bugs. In *Proceedings of the ACM International Conference on Object Oriented Programming Systems*